

# Dynamic Mixture of Experts: An Auto-Tuning Approach for Efficient Transformer Models

Published at ICLR 2025

Original Authors: Y. Guo et al.

Presented by: Fengnan Li

Duke B&B

Nov, 2025

# Outline

- 1 Introduction & Motivation
- 2 Methodology: Top-Any Gating
- 3 Methodology: Adaptive Architecture
- 4 Methodology: Loss Function
- 5 Experiments & Results
- 6 Conclusion

# The Bottleneck of Large Models

## The Trend:

- Transformers have scaled massively (Billions → Trillions of parameters).
- Dense scaling is computationally expensive (quadratic/linear cost increase).

## The Solution: Sparse Mixture of Experts (MoE)

- Replaces dense Feed-Forward Networks (FFN) with multiple "Experts".
- **Conditional Computation:** Only activates a subset of parameters per token.
- **Benefit:** High Capacity (Total Params) with Low Compute (Active Params).

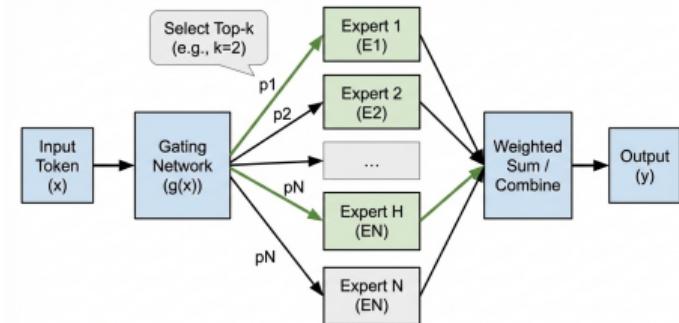


Diagram of Standard Sparse MoE Architecture

# The Problem: Hyperparameter Sensitivity

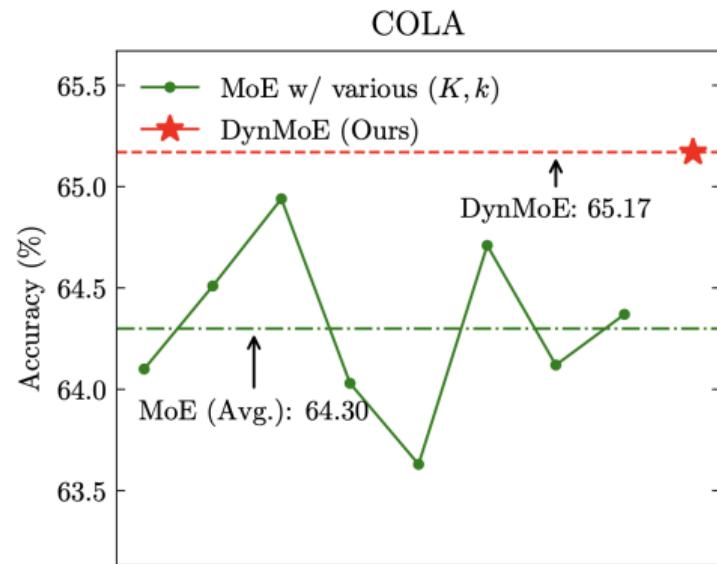
While MoE is efficient, training it is difficult due to sensitive hyperparameters.

## Critical Hyperparameters

- **$K$  (Total Experts):** How many experts should we have? (8, 16, 64?)
- **$k$  (Top- $k$  Activation):** How many experts should process each token? (1, 2, 4?)

## The Consequences:

- Performance fluctuates significantly (1%-3%) based on the choice of  $(K, k)$ .
- Finding the optimal pair requires expensive **Grid Search** or Ablation Studies.
- **Different tasks require different settings** (e.g., QNLI needs  $K = 16$ , RTE needs  $K = 8$ ).



(a) Performance Fluctuation Illustration

# The Proposal: Dynamic MoE (DynMoE)

**Core Question:** *Can we develop a training strategy that automatically determines both the number of experts and the activation count during training?*

## DynMoE Framework Overview:

- ① **Top-Any Gating:** A mechanism that allows tokens to autonomously decide how many experts to use (0 to N).
- ② **Adaptive Training Process:** A lifecycle management system that adds new experts when needed and removes idle ones.
- ③ **Auxiliary Loss:** A specialized loss to ensure the model remains sparse and numerically stable.

# 1. Top-Any Gating Mechanism

Standard MoE uses **Softmax** (Competition). DynMoE uses **Sigmoid** (Independence).

## Mechanism Details: Scoring & Discrete Decision

- ① **Cosine Similarity Scoring:** Calculate  $s(x)$  between input token  $x$  and expert weights  $W_g$ . Each expert has a corresponding column in  $W_g$ .

$$s(x) = \frac{\langle x, W_g \rangle}{\|x\| \cdot \|W_g\|}$$

- ② **Discrete Gating Decision**  $g(x)$ : Apply the **Sign Function** to convert soft scores into binary activations  $\{0, 1\}$ . This compares the sigmoid-scaled similarity score  $\sigma(s(x))$  against a learnable per-expert threshold  $\sigma(G)$ .

$$g(x) = \text{sign}(\sigma(s(x)) - \sigma(G))$$

- This gating allows a token to activate **any number** of experts (from 0 to  $N$ ).

# 1. Top-Any Gating Mechanism

**Continuing from Part 1:** After determining active experts using  $g(x)$ .

## Mechanism Details: Expert Output Aggregation

- ③ **Expert Output Aggregation:** Compute the simple average of outputs from all activated experts. Here,  $k = \sum g(x)$  is the count of active experts for the current token.

$$y = \frac{1}{k} \sum_{g(x_e) > 0} E_e(x)$$

## Crucial Design Questions:

- **Scale Mismatch:** Why use simple averaging ( $1/k$ ) instead of weighted sum (like standard MoE)? We'll discuss this soon.
- **Gradient Flow:** We used a non-differentiable 'sign' function in  $g(x)$ . How do we pass gradients to learn  $W_g$  and  $G$ ? (Next Slide!)

# Passing Gradients via Straight-Through Estimator

**The Problem:** The derivative of the ‘sign’ function used in the previous step is 0 almost everywhere. Gradients cannot flow back.

**The Solution (STE Trick):** We use the **Straight-Through Estimator** in the computational graph to bypass the sign function during backpropagation.

Let  $\text{soft} = \sigma(s(x)) - \sigma(G)$ .

## Computational Graph Implementation

$$g_{\text{final}}(x) = (g(x) - \text{soft}).\text{detach}() + \text{soft}$$

- **Forward Pass:** Terms cancel out. The result is exactly the discrete  $g(x)$ .
- **Backward Pass:** ‘`.detach()`’ blocks gradient flow in the first term.

$$\frac{\partial \mathcal{L}}{\partial g_{\text{final}}} \approx \frac{\partial \mathcal{L}}{\partial \text{soft}}$$

Gradients successfully flow back to update weights and thresholds.

# Output Aggregation: Why Simple Average?

Once experts are selected, how do we combine their outputs?

## Equation (6): Simple Averaging

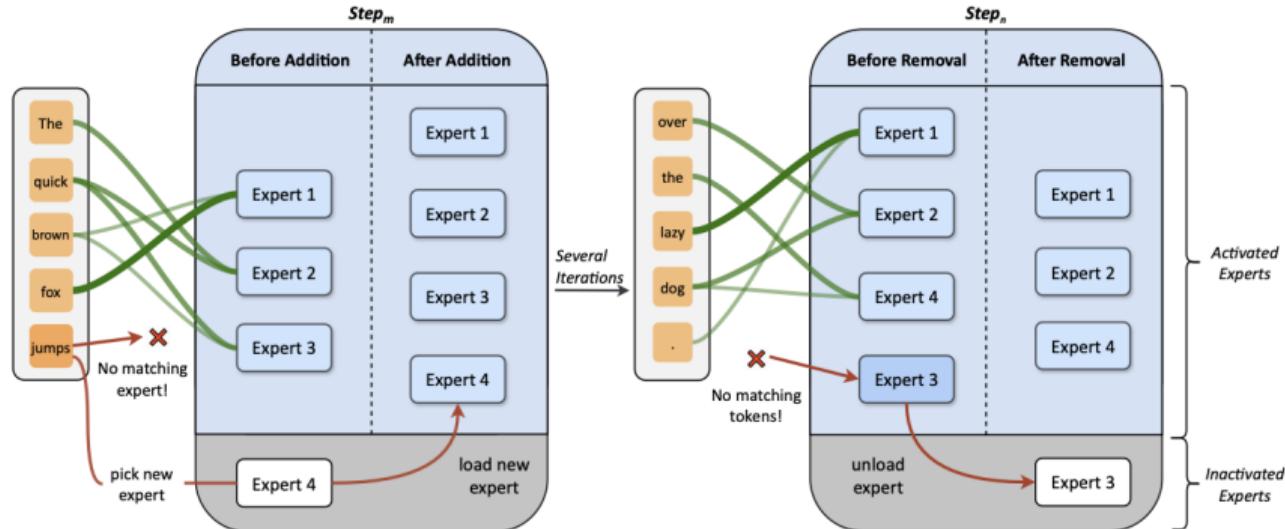
Let  $k = \sum g(x)$  be the number of activated experts.

$$y = \frac{1}{k} \sum_{g(x_e) > 0} E_e(x)$$

### Remark 3.1: The "Scale Mismatch" Problem

- Why not use scores as weights (like standard MoE)?
- In our "Independent Gating" approach, experts learn different thresholds.
- **Example:** Expert A might be activated with score 0.15 (Low Threshold), while Expert B is activated with 0.9 (High Threshold).
- Using scores as weights would unfairly penalize Expert A, even though both are valid experts.
- **Conclusion:** We ignore magnitude and treat all active experts equally.

## 2. Adaptive Training Process (Overview)



- **Timing:** The system checks and acts every **100-300 iterations** to stabilize the process.
- **Growth (Adding Experts):** If tokens are "homeless" ( $R_S \neq 0$ ), a new expert is **added** and initialized using the  $R_S$  direction.
- **Pruning (Removing Experts):** If an expert is completely idle ( $R_E = 0$ ), it is **removed** to maintain sparsity and efficiency.

# Phase 1: Recording "Homeless" Tokens

## The Problem of $k = 0$

Since thresholds are independent, it is possible that a token activates **zero** experts. This means the current experts are insufficient to cover this token's semantics.

## The Accumulator ( $R_S$ ):

- During the monitoring window, we identify tokens where  $\text{sum}(g(x)) = 0$ .
- We accumulate their embedding vectors into a buffer called  $R_S$ :

$$R_S = R_S + \sum_{i=1}^N \mathbb{1}_{k_i=0} \cdot x_i$$

- $R_S$  effectively represents the "average direction" of the unsolved problems.

## Phase 2: Adding New Experts

If  $R_S \neq 0$  (i.e., there are homeless tokens), we recruit a new expert.

### Initialization Strategy (Crucial)

We create Expert  $K + 1$ . How do we initialize it?

$$W_{g,K+1} = \frac{R_S}{\|R_S\|}, \quad G_{K+1} = 0$$

- **Weight Initialization:** We set the new expert's vector to point exactly in the direction of the homeless tokens ( $R_S$ ). This guarantees they will match next time.
- **Threshold Initialization:** We set  $G = 0$  (Sigmoid output 0.5). This is a "neutral" starting point, ensuring the expert is easily activated immediately.

# Adaptive Process: Removing Experts (Pruning)

The Adaptive Training Process also prunes underutilized experts to maintain efficiency.

## Mechanism: Pruning Idle Experts

- **Routing Record ( $R_E$ ):** During the monitoring window, the system records the activation count for every expert.
- **Condition:** We check if there is an expert  $e$  such that its recorded activation time is zero ( $R_E^e = 0$ ).
- **Action:** If the expert  $e$  has been completely idle, it is **removed** from the model structure to free up computational resources.

### 3. Sparse and Simple Gating Loss

Without constraints, Top-Any Gating might activate ALL experts to minimize loss. We need a regularization term.

$$\mathcal{L}_{aux} = \underbrace{\|W_g^T W_g - I_K\|_2}_{\text{Diversity Loss}} + \underbrace{\frac{1}{K} \sum_{e=1}^K \|w_{g,e}\|_2}_{\text{Simplicity Loss}}$$

- **Diversity Loss ( $O(W^4)$ ):**

- Forces  $W_g^T W_g \approx I$ .
- This pushes expert vectors to be **Orthogonal**.
- If vectors are orthogonal, a single token cannot be similar to all of them simultaneously → **Enforces Sparsity**.

- **Simplicity Loss ( $O(W^2)$ ):**

- Penalizes the norm of weight vectors (L2 Regularization).
- **Why needed?** Diversity loss (4th order) can cause exploding gradients. Simplicity loss provides a linear counter-force to stabilize values and prevent Sigmoid saturation.

# Experiment Overview: 4 Key Questions

The authors designed experiments to answer four specific questions:

## Performance

- Q1:** Can DynMoE achieve competitive performance vs. fixed settings?
- Q2:** Can it handle different modalities (Vision/Language)?

## Efficiency & Insight

- Q3:** Does it maintain sparsity (Efficiency)?
- Q4:** What insights does it offer for MoE design?

# Q1: Competitive Performance (GLUE Benchmark)

Baseline: Standard MoE with various grid-searched  $(K, k)$ .

## Observations:

- **Green Lines (Standard MoE):** Show high variance. Optimal  $K, k$  changes for every task.
- **Red Star (DynMoE):** Consistently performs above average or at the top.
- **Top-2 Accuracy Times:** DynMoE achieves the highest score (3.00), proving it is the most robust setting across tasks.

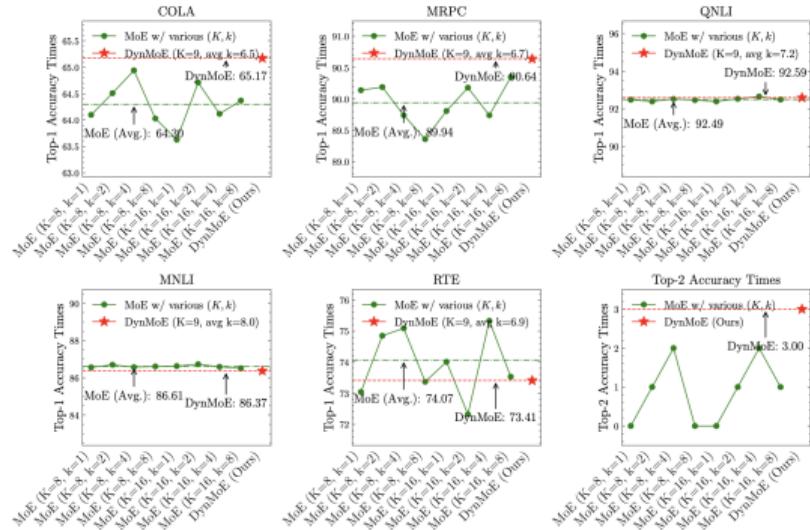


Figure 4: GLUE Results (Fluctuation vs. Robustness)

## Q2: Modalities & Scales

*DynMoE was tested on Vision (DomainBed) and Vision-Language (LLaVA) tasks.*

Task / Backbone	Baseline (MoE-LLaVA)	DynMoE (Ours)
<b>Vision (DomainBed)</b>	72.8 (avg)	72.4 (avg)
<b>VL (StableLM-1.6B)</b>	76.7	<b>77.4</b>
<b>VL (Qwen-1.8B)</b>	76.2	<b>76.4</b>
<b>VL (Phi-2-2.7B)</b>	77.6	<b>77.9</b>

**Table:** Performance Comparison on Vision-Language Tasks (VQA)

**Conclusion:** DynMoE consistently matches or outperforms carefully tuned baselines across different model scales and modalities.

# Q3: Sparsity & Efficiency (Does Auto-Tuning explode?)

"Does Auto-Tuning lead to parameter explosion?" No.

## Inference Efficiency

DynMoE automatically learns to activate **fewer** parameters than the fixed baseline ( $k = 2$ ).

- **StableLM:** Activated **1.75B** (DynMoE) vs 2.06B (Baseline).  
→ **~15% reduction in inference compute.**
- **Qwen:** 2.19B (DynMoE) vs 2.24B (Baseline).

## Training Efficiency

- Training FLOPs are comparable to Standard MoE.
- **Major Gain:** Saves the massive compute cost of "Ablation Studies" (running the training 8+ times to find the best  $K, k$ ).

# Q4: Insights for Architecture Design

The DYNMOE framework provided two key insights that can guide the future design of MoE models.

## Insight 1: Layer-wise Sparsity

- **Bottom Layers (Shallow):** Tokens tend to activate experts more uniformly (higher  $k$ ), reflecting the need for broader feature extraction.
- **Top Layers (Deep):** Tokens select fewer experts, often just one expert, as representations converge for final tasks.
- **Implication:** MoE structure may only be required for **bottom layers**. Top layers could remain dense, saving computation without performance loss.

## Insight 2: Emergence of Shared Experts

- **Observation:** Visualizing the learned thresholds ( $G$ ) showed that one expert per layer consistently learns a significantly **lower threshold**.
- **Pattern:** This low threshold makes that expert more readily activated by almost all tokens, acting as a "Generalist".
- **Validation:** This self-emergent behavior is consistent with the manual design choice of incorporating shared experts in models like DeepSeek-MoE.

# Summary

① **Problem:** DynMoE addresses the inefficiency of hyperparameter tuning in MoE training.

② **Method:**

- **Top-Any Gating** (Sigmoid + STE) allows flexible activation.
- **Adaptive Process** adds/removes experts based on real-time demand.

③ **Result:**

- Competitive performance with State-of-the-Art.
- Reduced inference costs (15% less params).
- Removed the need for grid search.

**Code available at:** <https://github.com/LINs-lab/DynMoE>

# Thank You!

Questions?