

MetaFormer is Actually What You Need for Vision

Yu, Weihao, Mi Luo, Pan Zhou, Chenyang Si, Yichen Zhou, Xinchao Wang, Jiashi Feng, and
Shuicheng Yan

Sea AI Lab & National University of Singapore

April 3, 2023

Presented by Boyao Li

- Introduction
- Background
- MetaFormer
- Experiments and Future Work

This paper abstracts transformers into a general architecture MetaFormer, and empirically demonstrates that the success of transformer/MLP-like models is largely attributed to the MetaFormer architecture rather than specific token mixers.

- PoolFormer: only employ a pooling operator as a weak token mixer for MetaFormer.
- Achieve competitive performance compared with the SOTA models using sophistic design of token mixers on multiple vision tasks.

Why matters?

- Inspire more future research dedicated to improving MetaFormer instead of the token mixer modules.
- PoolFormer could serve as a good start baseline for future MetaFormer architecture design.

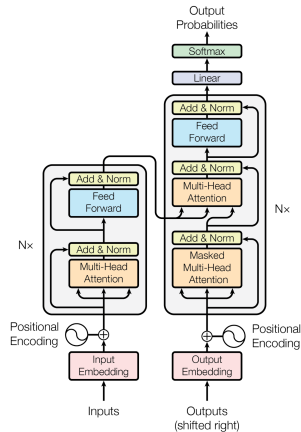
Background

- Transformer
- Vision Transformer (ViT)
- MLP-like Models

Transformer

Encoder and Decoder Stacks

- N (In this work $N = 6$) identical layers. Each layer has two sub-layers.
- Output for each sub-layer:
$$Y = \text{LayerNorm}(X + \text{SubLayer}(X))$$



Transformer

Scaled Dot-Product Attention


- Queries(d_k), keys(d_k) and values(d_v)
- $\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V$

Multi-Head Attention


- $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$,
where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $W_i^Q \in \mathbb{R}^{d_{\text{model}} \times d_k}, W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}, W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}, W_i^O \in \mathbb{R}^{hd_v \times d_{\text{model}}}$

In this work, $d_{\text{model}} = 512, h = 8, d_{d_v} = d_{\text{model}}/h = 64$.

Scaled Dot-Product Attention




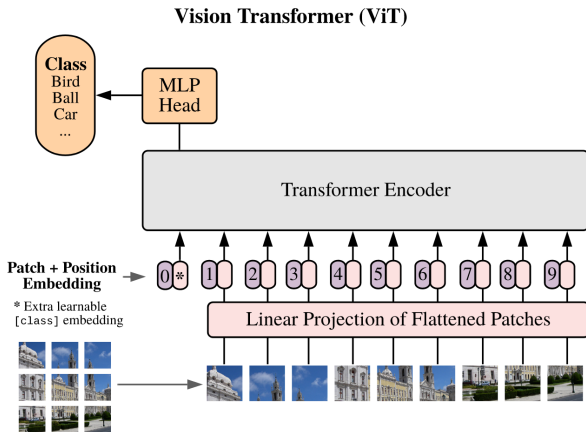
Multi-Head Attention



Vision Transformer(ViT)

- A pure transformer applied on sequences of image patches can perform very well on image classification tasks, and the reliance on CNN may not be necessary.
- Yield modest accuracies of a few percentage points below ResNets of comparable size when pre-trained on mid-sized datasets such as ImageNet-1K.
- Reach/beat SOTA models when trained on larger datasets(14M-300M images) such as ImageNet-21K(14M images with 21K classes) and JFK-300M dataset.

Vision Transformer(ViT)



Vision Transformer(ViT)

- Patch Embedding
 - Reshape a 2D (H, W) image $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ into a sequence of 2D (P, P) patches $\mathbf{x}_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$.
 - C is the number of the channels.
 - $N = HW/P^2$ is the number of the patches.
- Use standard learnable 1D position embeddings.
- Prepend a learnable embedding to the sequence of embedded patches $\mathbf{z}_0^0 = \mathbf{x}_{\text{class}}$ similar as BERT's CLS token. Only the class vector is used to predict the output.
- ViT has much less images-specific inductive bias than CNNs due to that self-attention layers are global.

Vision Transformer(ViT)

Related Work: Data-efficient Image Transformers (DeiT)

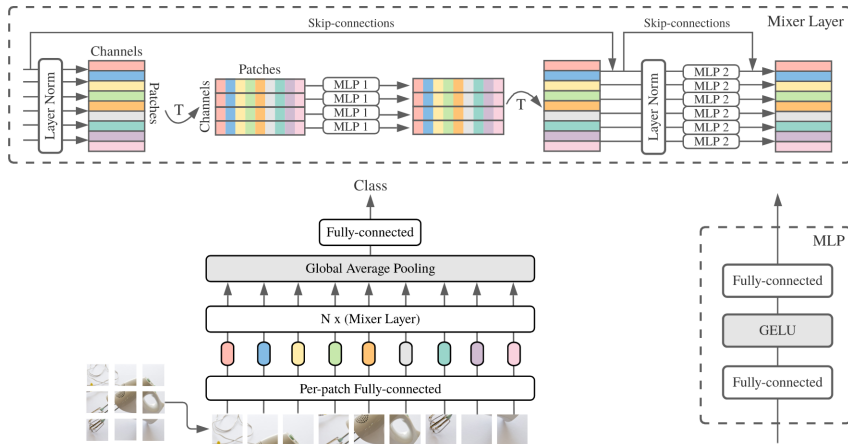
- Build upon ViT and produce competitive transformers by training on ImageNet only with less computing resource.
- New distillation strategy specific to transformers
 - Add a distillation token to the sequence of embedded patches to reproduce the label given by the "teacher" instead of true label on output of the network.
 - Both the class and distillation tokens are learned by backpropagation.

MLP-like Models

MLP-Mixer(and ResMLP)

- While convolutions and attentions are both sufficient for good performance, neither of them are necessary.
- Two types of MLP layers
 - token-mixing MLP: $\mathbb{R}^N \rightarrow \mathbb{R}^N$, applied across patches.
 - channel-mixing MLP: $\mathbb{R}^C \rightarrow \mathbb{R}^C$, applied independently to patches.
- Position embeddings are not included because token-mixing MLP are sensitive to the order of the input tokens.
- Attain competitive performance on image classification benchmarks when trained on large datasets or with modern regularization schemes.


MLP-Mixer




Related Work: Spatial-shift MLP

- MLP-Mixer cannot achieve as outstanding performance as its CNN/ViT counterparts.
- Replace token-mixing MLP to some spatial-shift operations for the communications between patches.


MetaFormer




MetaFormer
(General Arch.)



Transformer
(e.g. DeiT)



MLP-like model
(e.g. ResMLP)

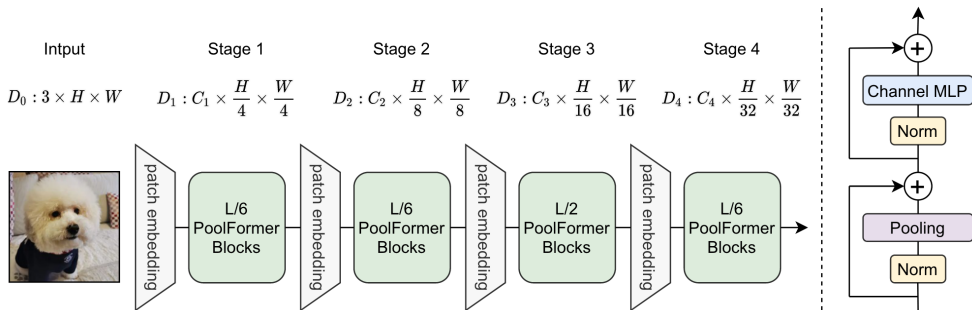


PoolFormer
(Ours)

- For an input image I , $X = \text{InputEmb}(I)$, $X \in \mathbb{R}^{N \times C}$
- $Y = \text{TokenMixer}(\text{Norm}(X)) + X$. $\text{Norm}(\cdot)$ denotes normalization such as Layer Normalization or Batch Normalization.
- $Z = \sigma(\text{Norm}(Y)W_1)W_2 + Y$ where $W_1 \in \mathbb{R}^{C \times rC}$, $W_2 \in \mathbb{R}^{rC \times C}$. r is the MLP (expansion) ratio, and $\sigma(\cdot)$ is a non-linear activation function such as GeLU or ReLU.

- Use (average) pooling as the token mixer.
- Assuming input $T \in \mathbb{R}^{C \times H \times W}$ is channel-first, the pooling operator is $T'_{:,i,j} = \left(\frac{1}{K \times K} \sum_{p,q=1}^K T'_{:,i+p-\frac{K+1}{2},i+q-\frac{K+1}{2}} \right) - T_{:,i,j}$, where K is the pooling size.
- The subtraction of the input itself is added due to that MetaFormer block already has a residual connection.

Overall Framework of PoolFormer



(a) Overall framework with L PoolFormer blocks

(b) PoolFormer block

Experiments: Model Configurations

Stage	#Tokens	Layer Specification		PoolFormer				
				S12	S24	S36	M36	M48
1	$\frac{H}{4} \times \frac{W}{4}$	Patch	Patch Size	7×7 , stride 4				
		Embedding	Embed. Dim.	64		96		
		PoolFormer Block	Pooling Size	3×3 , stride 1				
			MLP Ratio	4				
			# Block	2	4	6	6	8
2	$\frac{H}{8} \times \frac{W}{8}$	Patch	Patch Size	3×3 , stride 2				
		Embedding	Embed. Dim.	128		192		
		PoolFormer Block	Pooling Size	3×3 , stride 1				
			MLP Ratio	4				
			# Block	2	4	6	6	8
3	$\frac{H}{16} \times \frac{W}{16}$	Patch	Patch Size	3×3 , stride 2				
		Embedding	Embed. Dim.	320		384		
		PoolFormer Block	Pooling Size	3×3 , stride 1				
			MLP Ratio	4				
			# Block	6	12	18	18	24
4	$\frac{H}{32} \times \frac{W}{32}$	Patch	Patch Size	3×3 , stride 2				
		Embedding	Embed. Dim.	512		768		
		PoolFormer Block	Pooling Size	3×3 , stride 1				
			MLP Ratio	4				
			# Block	2	4	6	6	8
Parameters (M)				11.9	21.4	30.8	56.1	73.4
MACs (G)				1.9	3.5	5.1	9.0	11.8

Experiments: Image Classification

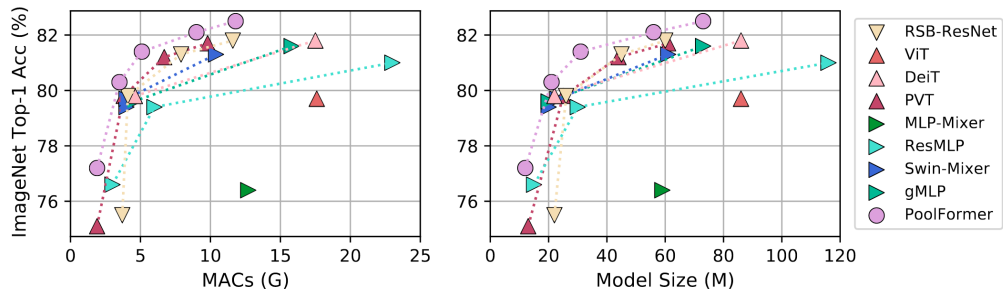
Setup

- Dataset: ImageNet-1K
- Data augmentation: MixUp, CutMix, CutOut and RandAugment
- Use Modified Layer Normalization (MLN) to compute the mean and variance along token and channel dimensions compared to only channel dimension in vanilla Layer Normalization.

Experiments: Image Classification

General Arch.	Token Mixer	Outcome Model	Image Size	Params (M)	MACs (G)	Top-1 (%)
Convolutional Neural Netowrks	—	▽ RSB-ResNet-18 [57]	224	12	1.8	70.6
		▽ RSB-ResNet-34 [57]	224	22	3.7	75.5
		▽ RSB-ResNet-50 [57]	224	26	4.1	79.8
		▽ RSB-ResNet-101 [57]	224	45	7.9	81.3
		▽ RSB-ResNet-152 [57]	224	60	11.6	81.8
MetaFormer	Attention	▲ ViT-B/16* [17]	224	86	17.6	79.7
		▲ ViT-L/16* [17]	224	307	63.6	76.1
		▲ DeiT-S [51]	224	22	4.6	79.8
		▲ DeiT-B [51]	224	86	17.5	81.8
		▲ PVT-Tiny [55]	224	13	1.9	75.1
		▲ PVT-Small [55]	224	25	3.8	79.8
		▲ PVT-Medium [55]	224	44	6.7	81.2
		▲ PVT-Large [55]	224	61	9.8	81.7
	Spatial MLP	► MLP-Mixer-B/16 [49]	224	59	12.7	76.4
		► ResMLP-S12 [50]	224	15	3.0	76.6
		► ResMLP-S24 [50]	224	30	6.0	79.4
		► ResMLP-B24 [50]	224	116	23.0	81.0
		► Swin-Mixer-T/D24 [35]	256	20	4.0	79.4
		► Swin-Mixer-T/D6 [35]	256	23	4.0	79.7
		► Swin-Mixer-B/D24 [35]	224	61	10.4	81.3
		► gMLP-S [34]	224	20	4.5	79.6
		► gMLP-B [34]	224	73	15.8	81.6
	Pooling	● PoolFormer-S12	224	12	1.9	77.2
		● PoolFormer-S24	224	21	3.5	80.3
		● PoolFormer-S36	224	31	5.1	81.4
		● PoolFormer-M36	224	56	9.0	82.1
		● PoolFormer-M48	224	73	11.8	82.5

Experiments: Image Classification



Experiments: Object Detection and Instance Segmentation

Backbone	RetinaNet 1×							Mask R-CNN 1×						
	Params (M)	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L	Params (M)	AP ^b	AP ₅₀ ^b	AP ₇₅ ^b	AP ^m	AP ₅₀ ^m	AP ₇₅ ^m
▼ ResNet-18 [23]	21.3	31.8	49.6	33.6	16.3	34.3	43.2	31.2	34.0	54.0	36.7	31.2	51.0	32.7
● PoolFormer-S12	21.7	36.2	56.2	38.2	20.8	39.1	48.0	31.6	37.3	59.0	40.1	34.6	55.8	36.9
▼ ResNet-50 [23]	37.7	36.3	55.3	38.6	19.3	40.0	48.8	44.2	38.0	58.6	41.4	34.4	55.1	36.7
● PoolFormer-S24	31.1	38.9	59.7	41.3	23.3	42.1	51.8	41.0	40.1	62.2	43.4	37.0	59.1	39.6
▼ ResNet-101 [23]	56.7	38.5	57.8	41.2	21.4	42.6	51.1	63.2	40.4	61.1	44.2	36.4	57.7	38.8
● PoolFormer-S36	40.6	39.5	60.5	41.8	22.5	42.9	52.4	50.5	41.0	63.1	44.8	37.7	60.1	40.0

Table 3. **Performance of object detection using RetinaNet, and object detection and instance segmentation using Mask R-CNN on COCO **val2017** [33].** 1× training schedule (*i.e.* 12 epochs) is used for training detection models. AP^b and AP^m represent bounding box AP and mask AP, respectively.

Experiments: Object Detection and Instance Segmentation

Backbone	RetinaNet 1×							Mask R-CNN 1×						
	Params (M)	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L	Params (M)	AP ^b	AP ₅₀ ^b	AP ₇₅ ^b	AP ^m	AP ₅₀ ^m	AP ₇₅ ^m
▼ ResNet-18 [23]	21.3	31.8	49.6	33.6	16.3	34.3	43.2	31.2	34.0	54.0	36.7	31.2	51.0	32.7
● PoolFormer-S12	21.7	36.2	56.2	38.2	20.8	39.1	48.0	31.6	37.3	59.0	40.1	34.6	55.8	36.9
▼ ResNet-50 [23]	37.7	36.3	55.3	38.6	19.3	40.0	48.8	44.2	38.0	58.6	41.4	34.4	55.1	36.7
● PoolFormer-S24	31.1	38.9	59.7	41.3	23.3	42.1	51.8	41.0	40.1	62.2	43.4	37.0	59.1	39.6
▼ ResNet-101 [23]	56.7	38.5	57.8	41.2	21.4	42.6	51.1	63.2	40.4	61.1	44.2	36.4	57.7	38.8
● PoolFormer-S36	40.6	39.5	60.5	41.8	22.5	42.9	52.4	50.5	41.0	63.1	44.8	37.7	60.1	40.0

Table 3. **Performance of object detection using RetinaNet, and object detection and instance segmentation using Mask R-CNN on COCO **val2017** [33].** 1× training schedule (*i.e.* 12 epochs) is used for training detection models. AP^b and AP^m represent bounding box AP and mask AP, respectively.

Experiments: Semantic segmentation

Backbone	Semantic FPN	
	Params (M)	mIoU (%)
▼ ResNet-18 [23]	15.5	32.9
▲ PVT-Tiny [55]	17.0	35.7
● PoolFormer-S12	15.7	37.2
▼ ResNet-50 [23]	28.5	36.7
▲ PVT-Small [55]	28.2	39.8
● PoolFormer-S24	23.2	40.3
▼ ResNet-101 [23]	47.5	38.8
▼ ResNeXt-101-32x4d [60]	47.1	39.7
▲ PVT-Medium [55]	48.0	41.6
● PoolFormer-S36	34.6	42.0
▲ PVT-Large [55]	65.1	42.1
● PoolFormer-M36	59.8	42.4
▼ ResNeXt-101-64x4d [60]	86.4	40.2
● PoolFormer-M48	77.1	42.7

Table 4. **Performance of Semantic segmentation on ADE20K [65] validation set.** All models are equipped with Semantic FPN [29].

Experiments: Ablation Studies

Ablation	Variant	Params (M)	MACs (G)	Top-1 (%)
Baseline	None (PoolFormer-S12)	11.9	1.9	77.2
Token mixers	Pooling \rightarrow Identity mapping	11.9	1.9	74.3
	Pooling \rightarrow Global random matrix* (extra 21M frozen parameters)	11.9	3.3	75.8
	Pooling \rightarrow Depthwise Convolution [9, 37]	11.9	1.9	78.1
	Pooling size 3 \rightarrow 5	11.9	1.9	77.2
	Pooling size 3 \rightarrow 7	11.9	1.9	77.1
	Pooling size 3 \rightarrow 9	11.9	1.9	76.8
Normalization	Modified Layer Normalization [†] \rightarrow Layer Normalization [1]	11.9	1.9	76.5
	Modified Layer Normalization [†] \rightarrow Batch Normalization [27]	11.9	1.9	76.4
	Modified Layer Normalization [†] \rightarrow None	11.9	1.9	46.1
Activation	GELU [24] \rightarrow ReLU [40]	11.9	1.9	76.4
	GELU \rightarrow SiLU [18]	11.9	1.9	77.2
Other components	Residual connection [24] \rightarrow None	11.9	1.9	0.1
	Channel MLP \rightarrow None	2.5	0.3	5.7
Hybrid Stages	[Pool, Pool, Pool, Pool] \rightarrow [Pool, Pool, Pool, Attention]	14.0	2.0	78.3
	[Pool, Pool, Pool, Pool] \rightarrow [Pool, Pool, Attention, Attention]	16.5	2.6	81.0
	[Pool, Pool, Pool, Pool] \rightarrow [Pool, Pool, Pool, SpatialFC]	11.9	1.9	77.5
	[Pool, Pool, Pool, Pool] \rightarrow [Pool, Pool, SpatialFC, SpatialFC]	12.2	1.9	77.9

Future Work

- MetaFormer Baselines for Vision
 - IdentityFormer, RandFormer
 - ConvFormer, CAFormer (Convs+Attentions)
- Work on NLP tasks?

References

- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... Polosukhin, I. (2017). Attention is all you need. *Advances in neural information processing systems*, 30.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., ... Houlsby, N. (2020). An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*.
- Touvron, H., Cord, M., Douze, M., Massa, F., Sablayrolles, A., Jégou, H. (2021, July). Training data-efficient image transformers distillation through attention. In *International conference on machine learning* (pp. 10347-10357). PMLR.
- Tolstikhin, I. O., Houlsby, N., Kolesnikov, A., Beyer, L., Zhai, X., Unterthiner, T., ... Dosovitskiy, A. (2021). Mlp-mixer: An all-mlp architecture for vision. *Advances in neural information processing systems*, 34, 24261-24272.
- Touvron, H., Bojanowski, P., Caron, M., Cord, M., El-Nouby, A., Grave, E., ... Jégou, H. (2022). Resmlp: Feedforward networks for image classification with data-efficient training. *IEEE Transactions on Pattern Analysis and Machine Intelligence*.
- Yu, T., Li, X., Cai, Y., Sun, M., Li, P. (2022). S2-mlp: Spatial-shift mlp architecture for vision. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision* (pp. 297-306).
- Yu, W., Si, C., Zhou, P., Luo, M., Zhou, Y., Feng, J., ... Wang, X. (2022). Metaformer baselines for vision. *arXiv preprint arXiv:2210.13452*.