

Progress Reward Model for Reinforcement Learning via Large Language Models

Xiuhui Zhang, Ning Gao, Xingyu Jiang, Yihui Chen,
Yuheng Pan, Mohan Zhang, Yue Deng

Beihang University
NeurIPS 2025

December 12, 2025

Presented by Liwen Sun

Overview

- 1 Motivation
- 2 Background
- 3 Method Overview
- 4 Subtask Decomposition
- 5 Progress Reward Model
- 6 Theoretical Guarantees
- 7 Policy Training
- 8 Experiments
- 9 Discussion

The Challenge: Long-Horizon RL Tasks

Traditional RL struggles with:

- **Sparse rewards:** Only feedback at task completion
- **Long horizons:** Complex tasks require many steps
- **Multi-stage structure:** Need to accomplish subtasks in sequence

Example (Robot Pick-and-Place Task)

- ① Move gripper to object
- ② Grasp object
- ③ Move to goal location
- ④ Release object

Sparse reward: +1 only when object reaches goal. All intermediate steps get 0.

Why is this hard? Without intermediate feedback, agent must explore blindly.

Existing Solutions and Their Limitations

Traditional Approaches

- **Hierarchical RL:** Decompose autonomously
 - Requires expert data
 - Hard to scale
- **Reward shaping:** Design dense rewards
 - Needs domain expertise
 - Task-specific engineering

LLM-Augmented Approaches

- **LLM as Planner:** High-level decomposition
 - Lacks low-level guidance
 - Needs pre-trained policies
- **LLM as Rewarder:** Generate reward code
 - Hard for complex tasks
 - Requires extensive search

Key Gap: Current methods focus on *either* planning *or* reward, not both!

LLM-Augmented RL: Two Paradigms

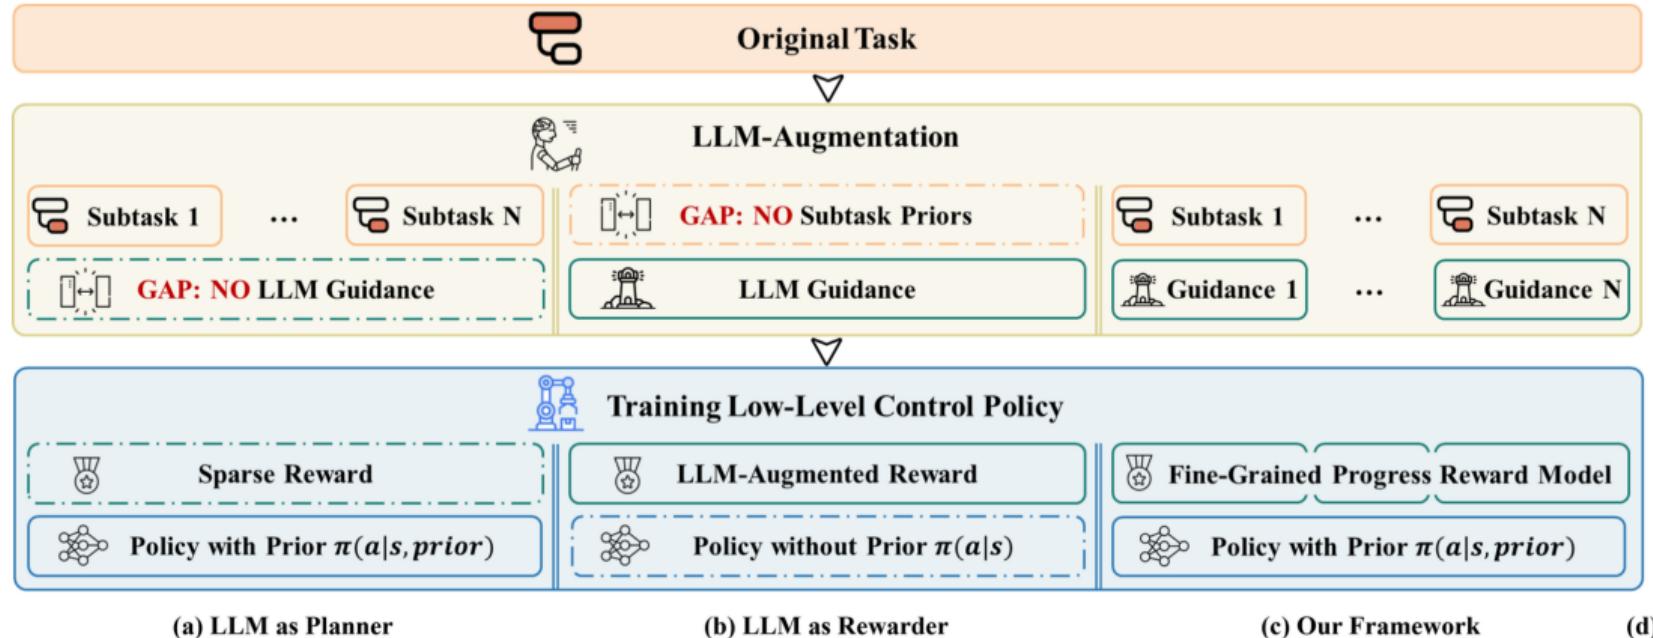


Figure: (a) LLM as Planner: Decompose task, but policy learns from sparse rewards. (b) LLM as Rewarder: Dense rewards, but no task decomposition. (c) PRM4RL: Combines both!

Our Contribution: PRM4RL

Progress Reward Model for RL (PRM4RL): Unified framework integrating planning & reward shaping

Key Idea

Use LLM to decompose task into subtasks, then construct a **progress function** that tracks execution.
Apply this as **potential-based reward shaping** for theoretically-grounded dense rewards.

Main Contributions:

- ① **High-level:** Subtask decomposition with automatic tracking (no repeated LLM calls)
- ② **Low-level:** Progress Reward Model with optimality & convergence guarantees
- ③ **Empirical:** State-of-the-art on MetaWorld and ManiSkill benchmarks
- ④ **Theoretical:** Formal proofs connecting progress to potential-based shaping

Background: Reinforcement Learning

Markov Decision Process (MDP): $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, T, R, \gamma, \rho_0 \rangle$

- \mathcal{S} : State space
- \mathcal{A} : Action space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$: Transition function
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$: Reward function
- $\gamma \in [0, 1]$: Discount factor
- ρ_0 : Initial state distribution

Goal: Find policy $\pi : \mathcal{S} \rightarrow \mathcal{A}$ maximizing expected return:

$$J(\pi) = \mathbb{E}_{\tau \sim p(\tau|\pi)} \left[\sum_{t=0}^T \gamma^t R(s_t, a_t) \right]$$

Q-function: Expected return from state-action pair:

$$Q^\pi(s, a) = \mathbb{E}_{\tau \sim p(\tau|\pi)} \left[\sum_{k=0}^T \gamma^k R(s_{t+k}, a_{t+k}) \mid s_t = s, a_t = a \right]$$

Background: Large Language Models for RL

Why use LLMs for RL?

- Pre-trained on diverse human knowledge
- Strong reasoning and planning capabilities
- Can generate code/structured outputs

Recent Approaches:

- **SayCan** (Brohan et al., 2023): LLM plans, selects from pre-trained policy library
- **Plan-Seq-Learn** (Dalal et al., 2024): LLM decomposes task, uses motion planning
- **ELLM** (Du et al., 2023): LLM generates plans, compute similarity as reward
- **Text2Reward** (Xie et al., 2024): LLM writes reward functions in Python
- **Eureka** (Ma et al., 2023): Evolutionary search over LLM-generated rewards

Gap: Planning methods lack dense rewards; reward methods lack decomposition

① Subtask Decomposition

- LLM breaks task into subtasks with "verb + noun" descriptions
- Generates determination function $\Psi(s)$ to identify current subtask

② Progress Function Construction

- For each subtask, LLM designs fine-grained progress metric
- Combine into overall progress $\Phi(s)$
- Use as potential function for reward shaping

③ Policy Training

- Augment state with subtask prior embedding
- Train with Progress Reward Model (PRM)
- Use standard RL algorithms (SAC/PPO)

Step 1: Prompting the LLM

Goal: Decompose complex task into manageable subtasks

Prompt Design (Pythonic representation):

- **Environment Info:** State space definition as Python classes
 - Robot attributes: end-effector position, joint states
 - Object attributes: position, orientation, initial states
 - Goal specification
- **Task Description:** Natural language specification
- **Instructions:** Chain-of-Thought reasoning guide

Output Format: Python code with comments for reasoning

- Comments (#) for intermediate thoughts
- Executable code for final outputs

Prompting Strategy: Overview

Goal: Generate code that tracks progress during task execution

Prompt Structure (following Text2Reward's Pythonic approach):

- ① **Basic Prompt:** Establishes role and objective
- ② **Environment Information:** Defines state space as Python classes
- ③ **Task Information:** Natural language task description
- ④ **Basic Instructions:** Code generation guidelines
- ⑤ **Reasoning Instructions:** Step-by-step procedure

Output: LLM generates `plan_list`, $\Psi(s)$, and $\Phi(s)$ in single call

Key Innovation: Program-aided Chain-of-Thought

- Reasoning in comments (#)
- Implementation in executable code

Prompting: Basic Prompt

Basic Prompt

You are an expert in robotics, reinforcement learning and code generation. We are going to use a robot arm to complete given tasks. Now I want you to help me write a python function named 'progress function' of reinforcement learning.

Purpose:

- Establishes LLM's role as expert
- Focuses on developing a `progress_function`
- Sets context: tracking progress to guide low-level policy learning

Prompting: Environment Information

Pythonic Class Representation:

Environment Classes

```
class BaseEnv(gym.Env):
    self.robot: Robot # the robot in the environment
    self.obj: RigidObject # the first object in the environment
    self.goal_position: np.ndarray[(3,)] # indicate the 3D position of the goal
    near_object = float(tcp_to_obj <= 0.3)
    success = float(obj_to_target <= 0.07)

class Robot:
    self.ee_position: np.ndarray[(3,)] # indicate the 3D position of effector
    self.hand_init_pos: np.ndarray[(3,)] # indicate the initial 3D position
    self.tcp = self.ee_position - [0, 0, 0.045] # tool center point

class RigidObject:
    self.position: np.ndarray[(3,)] # indicate the 3D position of object
    self.quaternion: np.ndarray[(4,)] # indicate the quaternion of object
    self.obj_init_pos: np.ndarray[(3,)] # indicate the initial 3D position
```

Benefits: Compact, informative, bootstraps Python code generation

Prompting: Task Information

Task Information

{ Fill in the task description here. }

Example: In pick-place, the robot need to pick up the object and move it to the goal position.

Purpose:

- Provides natural language specification of the goal
- User fills in specific task description
- Crucial for LLM to understand what to decompose

Other Examples:

- “Close a drawer by its handle”
- “Rotate the faucet handle counter-clockwise”
- “Press a button in y coordination”

Prompting: Basic Instructions

Basic Instruction

- ① You are allowed to use any existing python package if applicable. But only use these packages when it's really necessary.
- ② Do not invent any variable or attribute that is not given.
- ③ Think step by step, add comment for reasoning and thought when you write code.

Rationale:

- **Instruction 1:** Minimize dependencies, only import when needed (e.g., numpy)
- **Instruction 2:** Stay within defined environment - ensures code is executable
- **Instruction 3:** Explainability and debugging through comments

Reasoning Instructions(simplified)

Follow the steps below:

- a. Decompose the task and generate plan_list
- b. Write a 'determination_function' for determine which subtask we are in
- c. Write a 'progress_function'

Strategy: Step-by-step guidance for LLM

Step 2: LLM Generates Subtask List

Output 1: plan_list - ordered subtask descriptions

Example (Pick-Place Task)

```
plan_list = ['move to object', 'grasp object', 'move to goal', 'release object']
```

Design Choice: "verb + noun" format

- **Consistency:** Same patterns across different tasks
- **Generalization:** Transfers to unseen tasks with similar structure
- **Compositionality:** Natural building blocks for complex behaviors

Common patterns: "move to X", "grasp X", "open X", "close X", "push X", etc.

Step 3: Determination Function

Challenge: How to know which subtask the agent is currently executing?

Naive Solution: Call LLM at every timestep

- Too slow and expensive
- Used by some prior work (ELLM)

Our Solution: Generate determination function $\Psi(s)$ once

Determination Function

$$\Psi : \mathcal{S} \rightarrow \{0, 1, \dots, N - 1\}$$

Maps current state s to index of current subtask in `plan_list`.

Key Insight: With good understanding of state space, can determine subtask from state alone!

Example: Determination Function

Generated Python Code

```
def determination_function(env, plan_list):
    # Subtask 1: 'move to object'
    if np.linalg.norm(env.robot.tcp - env.obj.position) > 0.3:
        return 0 # Still moving towards object
    # Subtask 2: 'grasp object'
    elif env.obj.position[2] <= env.obj.obj_init_pos[2]:
        return 1 # Object not yet grasped (still at initial height)
    # Subtask 3: 'move to goal'
    elif np.linalg.norm(env.obj.position - env.goal_position) > 0.07:
        return 2 # Still moving towards goal
    # Subtask 4: 'release object'
    else:
        return 3 # Task completed
```

Logic: Each subtask has completion condition based on state features

Constructing the Progress Function

Goal: Fine-grained tracking of task execution progress

For each subtask i : Define sub-progress $\phi_i(s) \in [0, 1]$

- 0 = subtask just started
- 1 = subtask completed

Common sub-progress metrics:

- **Distance-based:** $\phi(s) = 1 - \frac{d_{\text{current}}}{d_{\text{initial}}}$
- **Height-based:** $\phi(s) = \frac{h_{\text{current}} - h_{\text{initial}}}{h_{\text{target}} - h_{\text{initial}}}$
- **Angle-based:** For rotation tasks

Example: Progress Function

Generated Python Code (Simplified)

```
def progress_function(env):
    subtask_idx = determination_function(env, plan_list)
    if subtask_idx == 0:  # 'move to object'
        init_dist = ||tcp_init - obj_init||
        curr_dist = ||tcp - obj||
        subprogress = 1 - curr_dist / init_dist
    elif subtask_idx == 1:  # 'grasp object'
        subprogress = (obj.z - obj_init.z) / 0.5
    elif subtask_idx == 2:  # 'move to goal'
        init_dist = ||obj_init - goal||
        curr_dist = ||obj - goal||
        subprogress = 1 - curr_dist / init_dist
    main_progress = subtask_idx + subprogress
    return main_progress, subtask_idx
```

From Progress to Reward

Key Idea: Use progress function $\Phi(s)$ as potential function!

Progress Reward Model (PRM)

$$R_t^{PRM} = \gamma \cdot \Phi(s_{t+1}) - \Phi(s_t) + I(s_{t+1}) \cdot r_{bonus}$$

Components:

- $\gamma \cdot \Phi(s_{t+1}) - \Phi(s_t)$: **Progress reward** (potential-based shaping)
 - γ : discount factor
 - Positive when making progress
 - Zero when stuck
 - Negative if regressing (rare)
- $I(s_{t+1}) \cdot r_{bonus}$: **Terminal reward** when task succeeds
 - Maintains original task objective
 - Typically $r_{bonus} = 10$ or 100

Result: Dense feedback at every step while preserving optimal policy!

Why Does This Work?

Intuition: Progress reward guides exploration toward completion

Example (Pick-Place Scenario)

- Agent moves gripper toward object: Φ increases \rightarrow positive reward
- Agent grasps object: Φ increases \rightarrow positive reward
- Agent moves object toward goal: Φ increases \rightarrow positive reward
- Agent moves away from goal: Φ decreases \rightarrow negative reward

Comparison to alternatives:

- **Sparse reward only:** No guidance until completion
- **Distance reward only:** May not capture multi-stage structure
- **Progress directly as reward:** Can lead to unintended behaviors (see ablation)
- **PRM (potential-based):** Theoretically grounded, guides efficiently

Theorem 1: Policy Invariance

Theorem (Policy Invariance)

Given MDP $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, T, R, \gamma, \rho_0 \rangle$ with sparse reward $R_t = I(s_{t+1}) \cdot r_{bonus}$, and potential-based shaping:

$$F = \gamma \cdot \Phi(s_{t+1}) - \Phi(s_t)$$

Let $\mathcal{M}' = \langle \mathcal{S}, \mathcal{A}, T, R + F, \gamma, \rho_0 \rangle$. Then every optimal policy in \mathcal{M}' is also optimal in \mathcal{M} , and vice versa.

What does this mean?

- PRM doesn't change what the optimal policy should do
- No "reward hacking" or "deceptive" behaviors
- Agent still maximizes task success, just learns faster

Proof sketch: Show $Q_{\mathcal{M}'}^*(s, a) = Q_{\mathcal{M}}^*(s, a) - \Phi(s)$, so $\arg \max_a$ unchanged.

Theorem 2: Convergence Efficiency

Theorem (Convergence Efficiency)

Consider two MDPs: \mathcal{M} with reward $R + F$ and Q-function Q , and \mathcal{M}' with reward R and Q-function Q' . If we initialize:

$$Q'_0(s, a) = Q_0(s, a) + \Phi(s)$$

Then for all $t \geq 0$:

$$Q'_t(s, a) = Q_t(s, a) + \Phi(s)$$

Practical implication:

- PRM \equiv good Q-value initialization
- Progress function provides prior knowledge: reduces initial error $|Q_0 - Q^*|$
- Accelerates convergence without changing optimum
- Shaping term acts as curriculum: guides exploration to high-value regions

Theoretical Summary

Corollary: Convergence Guarantee

Potential-based reward shaping preserves convergence guarantees of RL algorithms.

Why these results matter:

- ① **Safety**: Won't learn wrong behavior due to reward shaping
- ② **Efficiency**: Provably faster convergence than sparse rewards
- ③ **Generality**: Applies to any RL algorithm (Q-learning, policy gradient, etc.)
- ④ **Design principle**: Progress is natural choice for potential function

Contrast with LLM-reward baselines:

- Text2Reward, Eureka: No theoretical guarantees
- May require extensive tuning or human feedback
- PRM: One-shot generation with formal properties

Augmenting the MDP

Two augmentations to improve learning:

1. State Augmentation: Add subtask prior information

- Compute current subtask: $i = \Psi(s)$
- Get subtask description: $\text{desc} = \text{plan_list}[i]$
- Embed description: $\mathcal{P}(s) = \text{Encoder}(\text{desc})$
 - Uses SimCSE sentence encoder
 - Produces fixed-size vector embedding
- Augmented state: $s' = [s; \mathcal{P}(s)]$

2. Reward Augmentation: Use PRM reward

$$R_t^{PRM} = \gamma \cdot \Phi(s_{t+1}) - \Phi(s_t) + I(s_{t+1}) \cdot r_{bonus}$$

Augmented MDP:

$$\mathcal{M}' = \langle \mathcal{S} + \mathcal{P}, \mathcal{A}, T, R^{PRM}, \gamma, \rho_0 \rangle$$

Training Procedure

Algorithm 1 PRM4RL Training

```
1: Input: Environment, task description, RL algorithm
2: Offline Phase (once):
3:   Construct Pythonic prompt with environment info
4:   Call LLM to generate: plan_list,  $\Psi(s)$ ,  $\Phi(s)$ 
5: Training Loop:
6: for episode = 1 to  $N$  do
7:   Initialize state  $s_0$ 
8:   for timestep  $t = 0$  to  $T$  do
9:     Compute subtask prior:  $\mathcal{P}(s_t) = \text{Encoder}(\text{plan\_list}[\Psi(s_t)])$ 
10:    Augment state:  $s'_t = [s_t; \mathcal{P}(s_t)]$ 
11:    Select action:  $a_t \sim \pi(s'_t)$ 
12:    Execute action, observe  $s_{t+1}$ 
13:    Compute reward:  $r_t = \gamma \cdot \Phi(s_{t+1}) - \Phi(s_t) + I(s_{t+1}) \cdot r_{bonus}$ 
14:    Update policy with  $(s'_t, a_t, r_t, s'_{t+1})$ 
15:   end for
16: end for
```

Experimental Setup

Environments: MetaWorld

- 7-DOF Sawyer robot
- Tabletop manipulation
- 10 tasks evaluated
- Examples: button-press, door-close, pick-place

ManiSkill

- 7-DOF Franka Panda
- High-fidelity physics
- 5 tasks evaluated
- Examples: LiftCube, PickCube, LiftPegUpright

Baselines:

- **ELLM**: LLM planner + subtask priors, sparse/subtask rewards
- **Text2Reward (T2R)**: LLM generates dense reward function
- **Oracle**: Expert-designed dense rewards (upper bound)

Evaluation: Success rate averaged over 5 random seeds

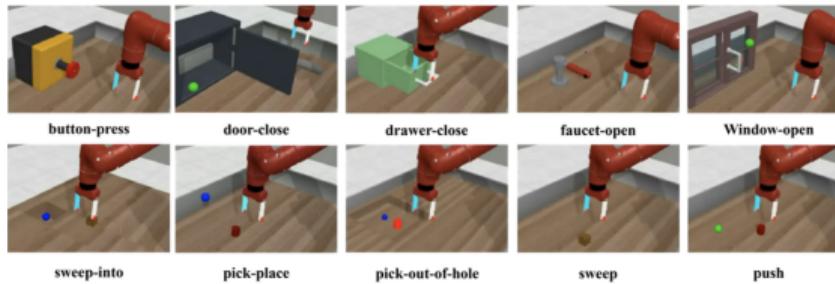


Figure 9: Tasks in Metaworld.

Table 3: Task list of Metaworld

Task	Task Description
button-press	Press a button in y coordination.
door-close	Close a door with a revolving joint by pushing the door's handle.
drawer-close	Close a drawer by its handle.
faucet-open	Rotate the handle counter-clockwise.
window-open	Push and open a sliding window by its handle.
sweep-into	Sweep a puck from the initial position into a hole.
pick-place	Pick up an object and move it to the goal location.
pick-out-of-hole	Pick an object out of a hole and move it to the goal location.
sweep	Sweep a puck off the table.
push	Push an object to the goal location.

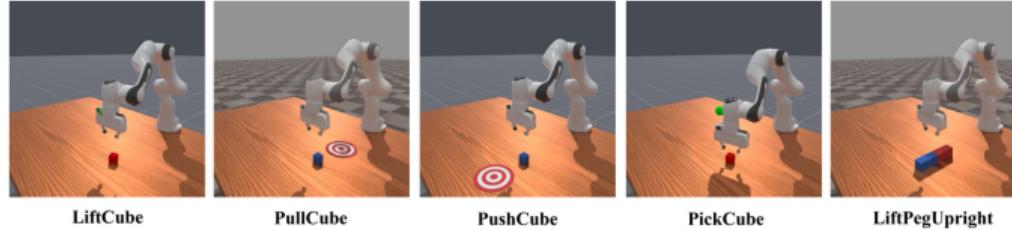


Figure 10: Tasks in Maniskill.

Table 4: Task list of Maniskill

Task	Task Description
LiftCube	Pick up cube A and lift it up by 0.2 meters.
PullCube	Pull a cube to the goal position.
PushCube	Push a cube to the goal position.
PickCube	Pick up cube A and move it to the 3D goal position.
LiftPegUpright	Move a peg laying on the table to any upright position on the table.

Main Results: Learning Curves

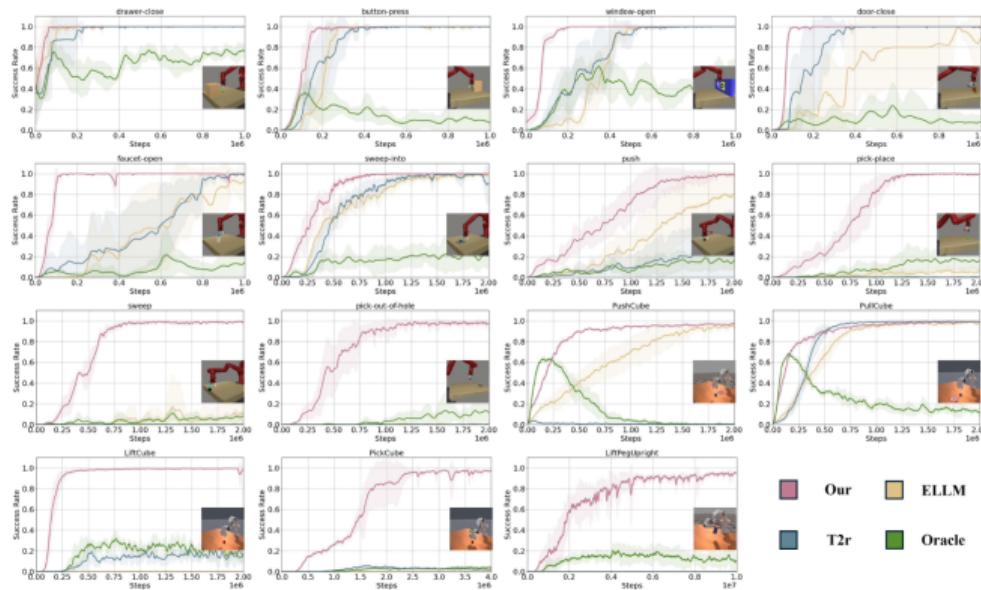


Figure 3: The learning curves of comparison methods on Metaworld and Maniskill measured by success rate. All experiments are conducted with 5 random seeds.

Figure: Training curves across tasks. PRM4RL (red) shows faster convergence and higher final performance than all baselines.

Main Results: Final Performance

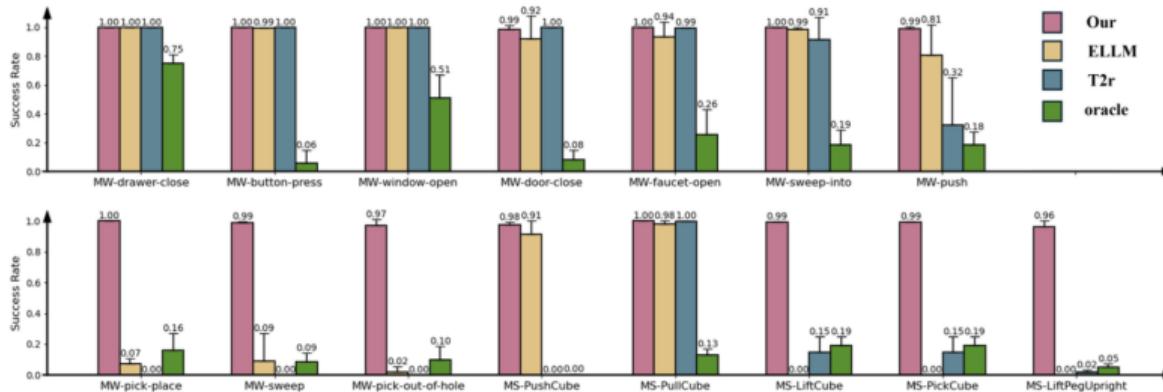


Figure 4: The evaluation results of comparison methods on Metaworld(With prefix ‘MW’) and Maniskill(with prefix ‘MS’) measured by success rate. All experiments are conducted with 5 random seeds.

Figure: Final success rates after training. PRM4RL achieves near-perfect performance on complex tasks where baselines fail.

Key Findings

1. Solves Complex Tasks:

- pick-out-of-hole: PRM4RL 100%, baselines <15%
- LiftPegUpright: PRM4RL 98%, baselines <10%
- Tasks with intricate multi-stage structure

2. Faster Convergence:

- Reaches high performance 2-3× faster than ELLM
- PRM rewards accelerate learning significantly

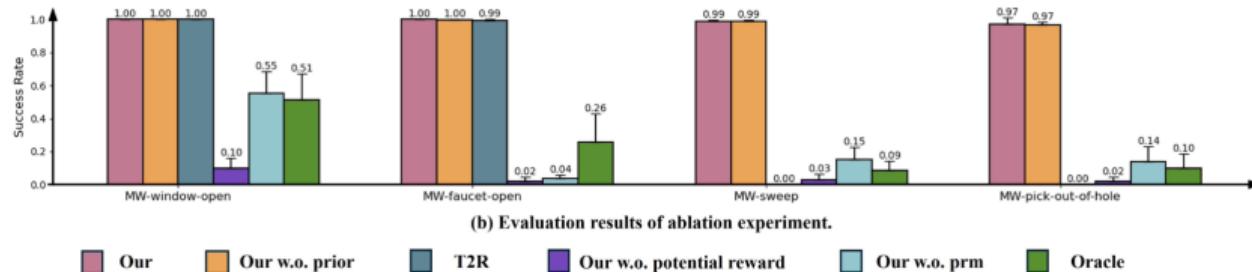
3. Outperforms Direct Reward Design:

- Beats Text2Reward on most tasks
- Subtask decomposition simplifies reward specification
- More reliable than evolutionary search

4. Approaches Oracle:

- Competitive with expert-designed rewards
- Sometimes even better (better exploration guidance)

Ablation Study: Both Components Matter



Variants tested:

- **w.o. PRM:** Subtask prior + oracle reward → Poor performance
- **w.o. prior:** PRM reward only → Slower convergence
- **w.o. potential:** Direct progress as reward → Unstable, degrades over time

Conclusion: Synergistic "1+1 >2" effect when combining both augmentations

Generalization to Unseen Tasks

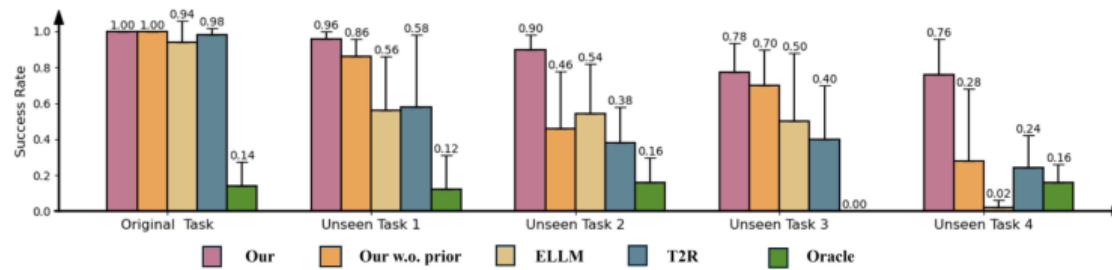


Figure 7: The evaluation results of generalization experiment. The results are averaged on 5 random seeds and 100 evaluations per seed.

Figure: Transfer performance: Train on one task, evaluate on related unseen tasks

Setup: Train on faucet-open, test on button-press, drawer-close, etc.

Results:

- PRM4RL shows smallest performance drop
- "verb + noun" pattern captures shared structure
- Subtask priors crucial for generalization (compare w.o. prior)

Strengths

① Theoretical Foundation

- Rare among LLM-RL papers to have formal guarantees
- Connects to established theory (potential-based shaping)

② Unified Framework

- Integrates planning and reward shaping naturally
- Each component addresses specific limitation

③ Efficiency

- Single LLM call vs. iterative search
- No per-timestep LLM invocation

④ Strong Empirical Results

- Consistent improvements across diverse tasks
- State-of-the-art on benchmarks

⑤ Thorough Ablations

- Clearly demonstrate necessity of both components
- Show importance of potential-based formulation

Limitations

1. LLM Hallucination Risk

- Method fails if LLM generates incorrect decomposition/progress
- No validation or error-handling discussed
- May need human verification in practice

2. Domain Specificity

- Only tested on robotics manipulation
- Requires structured state representations
- Unclear how to extend to other RL domains (games, dialogue, etc.)

3. Prompt Engineering

- Success depends on prompt quality
- Pythonic prompts require domain expertise
- Different environments need different templates

4. Limited Baseline Comparison

- Text2Reward uses mixed zero/few-shot settings
- Could potentially be stronger with more tuning

Open Questions

Practical Implementation:

- How to validate LLM-generated code automatically?
- What to do when determination function is inaccurate?
- How sensitive to choice of sentence encoder?

Generalization:

- Can this work for non-robotic RL domains?
- How to handle tasks without clear stage structure?
- What about partially observable environments?

Scaling:

- Does it work with smaller/cheaper LLMs?
- Can we learn to refine progress functions from experience?
- How to handle very long task horizons (>10 subtasks)?

Relevance to Healthcare

Clinical Decision-Making Applications:

1. Treatment Planning

- Decompose protocols: "stabilize patient" → "administer treatment" → "monitor recovery"
- Progress: Track patient state toward clinical milestones
- More interpretable than black-box RL

2. Discharge Recommendation

- Subtasks: "assess stability", "medication reconciliation", "arrange follow-up"
- Progress function: Multi-dimensional readiness score
- PRM framework for sequential decision support

3. Surgical Workflow

- Natural stage structure in procedures
- Progress tracking for real-time guidance
- Safety-critical: need theoretical guarantees!

Challenge: Medical LLMs need careful validation and domain-specific prompting

